

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: car_dataset = pd.read_csv(r"C:\Users\s323\Desktop\Gatherings\Data Science\Datasets\
```

```
In [3]: car_dataset.head()
```

```
Out[3]:
```

| | Car_Name | Year | Selling_Price | Present_Price | Kms_Driven | Fuel_Type | Seller_Type | Transmission |
|---|----------|------|---------------|---------------|------------|-----------|-------------|--------------|
| 0 | ritz | 2014 | 3.35 | 5.59 | 27000 | Petrol | Dealer | Manual |
| 1 | sx4 | 2013 | 4.75 | 9.54 | 43000 | Diesel | Dealer | Manual |
| 2 | ciaz | 2017 | 7.25 | 9.85 | 6900 | Petrol | Dealer | Manual |
| 3 | wagon r | 2011 | 2.85 | 4.15 | 5200 | Petrol | Dealer | Manual |
| 4 | swift | 2014 | 4.60 | 6.87 | 42450 | Diesel | Dealer | Manual |

```
In [4]: car_dataset.shape
```

```
Out[4]: (301, 9)
```

```
In [7]: car_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        301 non-null   object
1   Year            301 non-null   int64
2   Selling_Price   301 non-null   float64
3   Present_Price   301 non-null   float64
4   Kms_Driven      301 non-null   int64
5   Fuel_Type       301 non-null   object
6   Seller_Type     301 non-null   object
7   Transmission    301 non-null   object
8   Owner           301 non-null   int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

```
In [8]: car_dataset.describe()
```

Out[8]:

| | Year | Selling_Price | Present_Price | Kms_Driven | Owner |
|--------------|-------------|---------------|---------------|---------------|------------|
| count | 301.000000 | 301.000000 | 301.000000 | 301.000000 | 301.000000 |
| mean | 2013.627907 | 4.661296 | 7.628472 | 36947.205980 | 0.043189 |
| std | 2.891554 | 5.082812 | 8.644115 | 38886.883882 | 0.247915 |
| min | 2003.000000 | 0.100000 | 0.320000 | 500.000000 | 0.000000 |
| 25% | 2012.000000 | 0.900000 | 1.200000 | 15000.000000 | 0.000000 |
| 50% | 2014.000000 | 3.600000 | 6.400000 | 32000.000000 | 0.000000 |
| 75% | 2016.000000 | 6.000000 | 9.900000 | 48767.000000 | 0.000000 |
| max | 2018.000000 | 35.000000 | 92.600000 | 500000.000000 | 3.000000 |

In [9]: `car_dataset.columns`

Out[9]: Index(['Car_Name', 'Year', 'Selling_Price', 'Present_Price', 'Kms_Driven', 'Fuel_Type', 'Seller_Type', 'Transmission', 'Owner'], dtype='object')

In [10]: `car_dataset.isnull().sum()`

Out[10]: Car_Name 0
Year 0
Selling_Price 0
Present_Price 0
Kms_Driven 0
Fuel_Type 0
Seller_Type 0
Transmission 0
Owner 0
dtype: int64

In [12]: *# Checking the distribution of cateofrical data*
`print(car_dataset.Fuel_Type.value_counts())`
`print(car_dataset.Seller_Type.value_counts())`
`print(car_dataset.Transmission.value_counts())`

Petrol 239
Diesel 60
CNG 2
Name: Fuel_Type, dtype: int64
Dealer 195
Individual 106
Name: Seller_Type, dtype: int64
Manual 261
Automatic 40
Name: Transmission, dtype: int64

Label Encoding

In [17]: *# changing text/cateogorical value to numerical value*
Fuel_Type - Petrol is 0, Diesel is 1, CNG is 2
`car_dataset.replace({"Fuel_Type":{"Petrol":0,"Diesel":1,"CNG":2}},inplace=True)`
Seller_Type- Dealer is 0, Individual is 1
`car_dataset.replace({"Seller_Type":{"Dealer":0,"Individual":1}},inplace=True)`
Transmission- Manual is 0, Automatic is 1
`car_dataset.replace({"Transmission":{"Manual":0,"Automatic":1}},inplace=True)`

```
In [18]: car_dataset.head()
```

```
Out[18]:
```

| | Car_Name | Year | Selling_Price | Present_Price | Kms_Driven | Fuel_Type | Seller_Type | Transmission |
|---|----------|------|---------------|---------------|------------|-----------|-------------|--------------|
| 0 | ritz | 2014 | 3.35 | 5.59 | 27000 | 0 | 0 | 0 |
| 1 | sx4 | 2013 | 4.75 | 9.54 | 43000 | 1 | 0 | 0 |
| 2 | ciaz | 2017 | 7.25 | 9.85 | 6900 | 0 | 0 | 0 |
| 3 | wagon r | 2011 | 2.85 | 4.15 | 5200 | 0 | 0 | 0 |
| 4 | swift | 2014 | 4.60 | 6.87 | 42450 | 1 | 0 | 0 |

Splitting the data and Target

```
In [20]: X = car_dataset.drop(["Car_Name", "Selling_Price"], axis=1)  
Y = car_dataset["Selling_Price"]
```

```
In [21]: print (X)  
print (Y)
```

| | Year | Present_Price | Kms_Driven | Fuel_Type | Seller_Type | Transmission | \ |
|-----|------|---------------|------------|-----------|-------------|--------------|---|
| 0 | 2014 | 5.59 | 27000 | 0 | 0 | 0 | |
| 1 | 2013 | 9.54 | 43000 | 1 | 0 | 0 | |
| 2 | 2017 | 9.85 | 6900 | 0 | 0 | 0 | |
| 3 | 2011 | 4.15 | 5200 | 0 | 0 | 0 | |
| 4 | 2014 | 6.87 | 42450 | 1 | 0 | 0 | |
| .. | ... | ... | ... | ... | ... | ... | |
| 296 | 2016 | 11.60 | 33988 | 1 | 0 | 0 | |
| 297 | 2015 | 5.90 | 60000 | 0 | 0 | 0 | |
| 298 | 2009 | 11.00 | 87934 | 0 | 0 | 0 | |
| 299 | 2017 | 12.50 | 9000 | 1 | 0 | 0 | |
| 300 | 2016 | 5.90 | 5464 | 0 | 0 | 0 | |

| | Owner |
|-----|-------|
| 0 | 0 |
| 1 | 0 |
| 2 | 0 |
| 3 | 0 |
| 4 | 0 |
| .. | ... |
| 296 | 0 |
| 297 | 0 |
| 298 | 0 |
| 299 | 0 |
| 300 | 0 |

[301 rows x 7 columns]

| | |
|-----|-------|
| 0 | 3.35 |
| 1 | 4.75 |
| 2 | 7.25 |
| 3 | 2.85 |
| 4 | 4.60 |
| .. | ... |
| 296 | 9.50 |
| 297 | 4.00 |
| 298 | 3.35 |
| 299 | 11.50 |
| 300 | 5.30 |

Name: Selling_Price, Length: 301, dtype: float64

Splitting Training and Test data

```
In [22]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(X,Y,test_size=.1,random_state=2)
```

Model Training

```
In [23]: from sklearn.linear_model import LinearRegression
model = LinearRegression()
```

```
In [24]: model.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

Model Evaluation

```
In [31]: from sklearn.metrics import accuracy_score
```

```
In [32]: train_data_predictions = model.predict(x_train)
```


```
In [33]: from sklearn import metrics
```

```
In [40]: error_score = metrics.r2_score(y_train, train_data_predictions)
print("R squared Error : ", error_score)
```

R squared Error : 0.87994516604937

Visualize for accuracy - No accuracy score in Regression, it is only used for classification

```
In [44]: # visualize for actual price and predicted price
plt.figure(figsize=(5,5))
plt.scatter(y_train, train_data_predictions)
plt.xlabel("Actual prices")
plt.ylabel("Predicted prices")
plt.show()
```




```
In [45]: test_data_predictions = model.predict(x_test)
```

```
In [46]: error_score = metrics.r2_score(y_test, test_data_predictions)
print("R squared Error : ", error_score)
```

R squared Error : 0.8365766715025409

```
In [48]: # visualize for actual price and predicted price
plt.figure(figsize=(5,5))
plt.scatter(y_test, test_data_predictions)
plt.xlabel("Actual prices")
plt.ylabel("Predicted prices")
plt.show()
```



Lasso Regression

```
In [49]: from sklearn.linear_model import Lasso
```

```
In [51]: Lasso_model = Lasso()
```

```
In [52]: Lasso_model.fit(x_train,y_train)
```

```
Out[52]: Lasso()
```

Model Evaluation


```
In [53]: training_data_prediction = Lasso_model.predict(x_train)
```

```
In [54]: error_score = metrics.r2_score(y_train, training_data_prediction)
print("R squared Error : ", error_score)
```

```
R squared Error : 0.8427856123435794
```

Visualize the actual prices and Predicted prices

```
In [55]: plt.scatter(y_train, training_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")
plt.show()
```




```
In [58]: test_data_prediction = Lasso_model.predict(x_test)
```

```
In [61]: prinerror_score = metrics.r2_score(y_test, test_data_prediction)
("R squared Error : ", error_score)
```

```
Out[61]: ('R squared Error : ', 0.8427856123435794)
```

```
In [62]: plt.scatter(y_test, test_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")
plt.show()
```



```
In [ ]:
```